

Lab 9

```
## Intro to Data Science - Lab 9
# IST687 Section M002
# Professor Anderson
# Enter your name here: Chaithra Koppam Cheluvaiiah
# 1. I did this homework by myself, with help from the book and the professor.
```

```
#installing required packages
#install.packages("kernlab")
#install.packages("caret")
```

```
# importing the packages
library(kernlab)
library(caret)
```

```
## Loading required package: ggplot2
```

```
##
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:kernlab':
##
##      alpha
```

```
## Loading required package: lattice
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v tibble  3.1.5      v dplyr   1.0.7
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.0.2      v forcats 0.5.1
## v purrr   0.3.4
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x ggplot2::alpha() masks kernlab::alpha()
## x purrr::cross()   masks kernlab::cross()
## x dplyr::filter()  masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x purrr::lift()    masks caret::lift()
```

```
# getting GermanCredit data
data("GermanCredit")
subCredit <- GermanCredit[,1:10] # selecting only first 10 columns
glimpse(subCredit) # exploring the data
```

```
## Rows: 1,000
## Columns: 10
## $ Duration      <int> 6, 48, 12, 42, 24, 36, 24, 36, 12, 30, 12, 4~
## $ Amount        <int> 1169, 5951, 2096, 7882, 4870, 9055, 2835, 69~
## $ InstallmentRatePercentage <int> 4, 2, 2, 2, 3, 2, 3, 2, 2, 4, 3, 3, 1, 4, 2,~
## $ ResidenceDuration <int> 4, 2, 3, 4, 4, 4, 4, 2, 4, 2, 1, 4, 1, 4, 4,~
## $ Age           <int> 67, 22, 49, 45, 53, 35, 53, 35, 61, 28, 25, ~
## $ NumberExistingCredits <int> 2, 1, 1, 1, 2, 1, 1, 1, 1, 2, 1, 1, 1, 2, 1,~
## $ NumberPeopleMaintenance <int> 1, 1, 2, 2, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1,~
## $ Telephone      <dbl> 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1,~
## $ ForeignWorker  <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,~
## $ Class          <fct> Good, Bad, Good, Good, Bad, Good, Good, Good, Good~
```

1. Examine the data structure that `str()` reveals. Also use the `help()` function to learn more about the GermanCredit data set. Summarize what you see in a comment.

```
help(GermanCredit)
str(subCredit)
```

```
## 'data.frame': 1000 obs. of 10 variables:
## $ Duration      : int  6 48 12 42 24 36 24 36 12 30 ...
## $ Amount        : int  1169 5951 2096 7882 4870 9055 2835 6948 3059 5234 ...
## $ InstallmentRatePercentage: int  4 2 2 2 3 2 3 2 2 4 ...
## $ ResidenceDuration : int  4 2 3 4 4 4 4 2 4 2 ...
## $ Age           : int  67 22 49 45 53 35 53 35 61 28 ...
## $ NumberExistingCredits : int  2 1 1 1 2 1 1 1 1 2 ...
## $ NumberPeopleMaintenance : int  1 1 2 2 2 2 1 1 1 1 ...
## $ Telephone      : num  0 1 1 1 1 0 1 0 1 1 ...
## $ ForeignWorker  : num  1 1 1 1 1 1 1 1 1 1 ...
## $ Class          : Factor w/ 2 levels "Bad","Good": 2 1 2 2 1 2 2 2 2 1 ...
```

```
summary(subCredit)
```


```
##      Duration      Amount      InstallmentRatePercentage      ResidenceDuration
## Min.   : 4.0      Min.   : 250      Min.   :1.000      Min.   :1.000
## 1st Qu.:12.0      1st Qu.: 1366      1st Qu.:2.000      1st Qu.:2.000
## Median :18.0      Median : 2320      Median :3.000      Median :3.000
## Mean   :20.9      Mean   : 3271      Mean   :2.973      Mean   :2.845
## 3rd Qu.:24.0      3rd Qu.: 3972      3rd Qu.:4.000      3rd Qu.:4.000
## Max.   :72.0      Max.   :18424      Max.   :4.000      Max.   :4.000
##      Age      NumberExistingCredits      NumberPeopleMaintenance      Telephone
## Min.   :19.00      Min.   :1.000      Min.   :1.000      Min.   :0.000
## 1st Qu.:27.00      1st Qu.:1.000      1st Qu.:1.000      1st Qu.:0.000
## Median :33.00      Median :1.000      Median :1.000      Median :1.000
## Mean   :35.55      Mean   :1.407      Mean   :1.155      Mean   :0.596
## 3rd Qu.:42.00      3rd Qu.:2.000      3rd Qu.:1.000      3rd Qu.:1.000
```

```
## Max. :75.00 Max. :4.000 Max. :2.000 Max. :1.000
## ForeignWorker Class
## Min. :0.000 Bad :300
## 1st Qu.:1.000 Good:700
## Median :1.000
## Mean :0.963
## 3rd Qu.:1.000
## Max. :1.000
```

```
table(subCredit$Class)
```

```
##
## Bad Good
## 300 700
```

```
plot(subCredit$Class)
```



```
# Summarize what you see in a comment
# GermanCredit data set has 1000 rows and 10 columns
# 1) class represents credit worthiness of the customer and it is a dichotomous variable
# having only two possible values-bad and good
# 2) all other variables except class are quantitative variables
# 3) There are 300 customers classified as bad and 700 customers as good
```

2. Use the `createDataPartition()` function to generate a list of cases to include in the training data. This function is conveniently provided by `caret` and allows one to directly control the number of training cases. It also ensures that the training cases are balanced with respect to the outcome variable. Try this: `trainList <- createDataPartition(y=subCredit$Class,p=.40,list=FALSE)`

```
set.seed(111)
# randomly sample for training dataset elements
# 40% of data from subCredit = 400 cases will used for building the model
trainList <- createDataPartition(y=subCredit$Class,p=.40,list=FALSE)
```

3. Examine the contents of `trainList` to make sure that it is a list of case numbers. With `p=0.40`, it should have 400 case numbers in it.

```
trainList[1:20] # first 20 case numbers
```

```
## [1] 4 5 7 8 13 14 15 17 18 20 24 27 31 35 41 43 45 46 50 53
```

```
length(trainList)
```

```
## [1] 400
```

4. What is `trainList`? What do the elements in `trainList` represent? Which attribute is balanced in the `trainList` dataset?

```
# What is trainList?
# trainList is a vector consisting of case numbers that will used for building the SVM model

# What do the elements in trainList represent?
# case numbers/row indices from subCredit dataframe

# Which attribute is balanced in the trainList dataset?
# Class attribute
```

5. Use `trainList` and the square brackets notation to create a training data set called “`trainSet`” from the `subCredit` data frame. Look at the structure of `trainSet` to make sure it has all of the same variables as `subCredit`. The `trainSet` structure should be a data frame with 400 rows and 10 columns.

```
trainSet <- subCredit[trainList,]
head(trainSet)
```

```
##      Duration Amount InstallmentRatePercentage ResidenceDuration Age
## 4          42   7882                   2              4      45
## 5          24   4870                   3              4      53
## 7          24   2835                   3              4      53
## 8          36   6948                   2              2      35
## 13         12   1567                   1              1      22
## 14         24   1199                   4              4      60
##      NumberExistingCredits NumberPeopleMaintenance Telephone ForeignWorker Class
## 4                        1                        2          1          1   Good
## 5                        2                        2          1          1   Bad
```

```
## 7          1          1          1          1 Good
## 8          1          1          0          1 Good
## 13         1          1          0          1 Good
## 14         2          1          1          1 Bad
```

```
dim(trainSet)
```

```
## [1] 400 10
```


6. Use trainList and the square brackets notation to create a testing data set called “testSet” from the subCredit data frame. The testSet structure should be a data frame with 600 rows and 10 columns and should be a completely different set of cases than trainSet.

```
testSet <- subCredit[-trainList,]
dim(testSet)
```

```
## [1] 600 10
```

7. Create and interpret boxplots of all the predictor variables in relation to the outcome variable (Class).

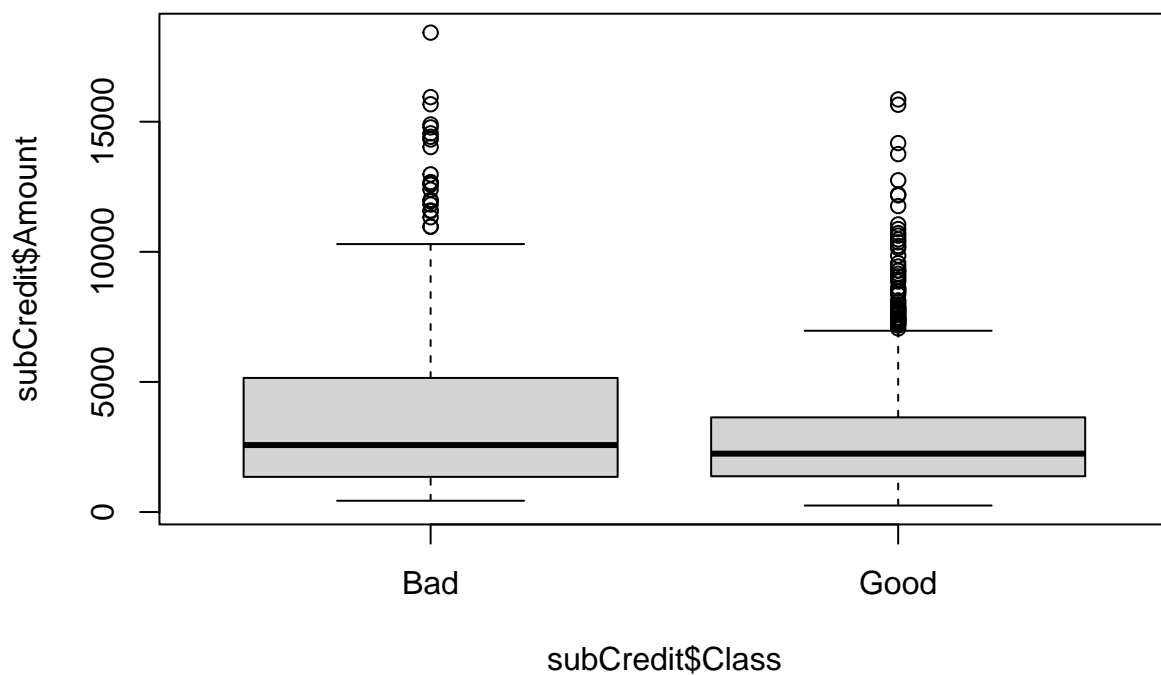
```
boxplot(subCredit$Duration~subCredit$Class,data=subCredit)
```



```
# Duration with class value Bad: data is not spread out. from the plot, we can  
# see first and second quartile are close to median with no outliers. This implies  
# most of the data is around mean/median
```

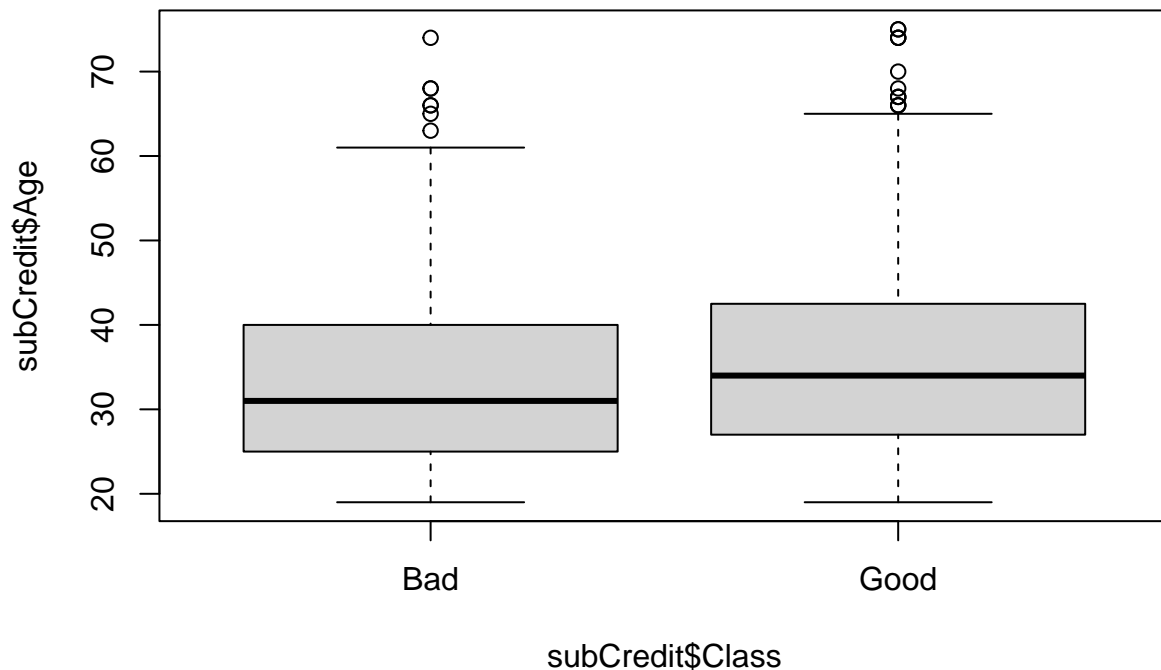
```
# Duration with class value Good: data is not spread out. from the plot, we can  
# see first and second quartile are close to median but it has many outliers  
# this results in right skewed distribution
```

```
boxplot(subCredit$Amount~subCredit$Class,data=subCredit)
```



```
# Amount variable: from the plot we can notice that there are many outliers in  
# the data resulting right skewed distribution
```

```
boxplot(subCredit$Age~subCredit$Class,data=subCredit)
```



*# Age Variable is also having many outliers in the data resulting in right skewed
distribution*

*# we need to handle outliers in order to get more accurate model while using these
features in the model building*

8. Train a support vector machine with the `ksvm()` function from the `kernlab` package. Make sure that you have installed and libraries the `kernlab` package. Have the cost be 5, and have `ksvm` do 3 cross validations (hint: try `prob.model = TRUE`)

```
?ksvm
svmModel <- ksvm(Class~.,data=trainSet,C=5,cross=3,prob.model=TRUE)
```

9. Examine the `ksvm` output object. In particular, look at the cross-validation error for an initial indication of model quality. Add a comment that gives your opinion on whether this is a good model.

```
svmModel

## Support Vector Machine object of class "ksvm"
##
## SV type: C-svc (classification)
## parameter : cost C = 5
##
## Gaussian Radial Basis kernel function.
```

```
## Hyperparameter : sigma = 0.124568955390236
##
## Number of Support Vectors : 251
##
## Objective Function Value : -846.3583
## Training error : 0.155
## Cross validation error : 0.279954
## Probability model included.
```

```
# cross validation error is approx 28%
# This indicates overall prediction error of all the 3 folds (cross=3)

# also, training error is 15.5%
# considering both the errors, it seems to be bad model as accuracy is
# is only ~56.5% (100% - (cross validation error%+training error%)) on training data set
```

10. Predict the training cases using the predict command

```
preOut <- predict(svmModel,newdata=testSet, type="response")
preOut[1:20]
```

```
## [1] Good Bad Good Good Good Good Good Good Bad Good Good Good Good Good Good Good
## [16] Good Good Bad Good Good
## Levels: Bad Good
```

11. Examine the predicted out object with str(). Then, calculate a confusion matrix using the table function.

```
str(preOut) # preOut is a vector of factor variable with two levels-Bad/Good
```

```
## Factor w/ 2 levels "Bad","Good": 2 1 2 2 2 2 2 1 2 2 ...
```

```
table(preOut,testSet$Class) #creating confusion matrix
```

```
##
## preOut Bad Good
## Bad 49 52
## Good 131 368
```

12. Interpret the confusion matrix and in particular calculate the overall accuracy of the model. The diag() command can be applied to the results of the table command you ran in the previous step. You can also use sum() to get the total of all four cells.

```
accuracy <- sum(diag(table(preOut,testSet$Class)))/sum(table(preOut,testSet$Class))
accuracy
```

```
## [1] 0.695
```



```
# Interpret the confusion matrix  
# diagonal elements indicate the right prediction of both values of Class  
# off-diagonal elements are the wrong predictions  
# so, accuracy = sum of all the right predictions/total predictions
```

13. Check you calculation with confusionMatrix() function in the caret package.

```
confusionMatrix(preOut,testSet$Class)
```

```
## Confusion Matrix and Statistics  
##  
##           Reference  
## Prediction Bad Good  
##      Bad    49   52  
##      Good  131  368  
##  
##           Accuracy : 0.695  
##           95% CI : (0.6564, 0.7316)  
##      No Information Rate : 0.7  
##      P-Value [Acc > NIR] : 0.6244  
##  
##           Kappa : 0.1697  
##  
##  McNemar's Test P-Value : 8.121e-09  
##  
##           Sensitivity : 0.27222  
##           Specificity : 0.87619  
##           Pos Pred Value : 0.48515  
##           Neg Pred Value : 0.73747  
##           Prevalence : 0.30000  
##           Detection Rate : 0.08167  
##      Detection Prevalence : 0.16833  
##           Balanced Accuracy : 0.57421  
##  
##           'Positive' Class : Bad  
##
```

```
# accuracy=69.5% is same as that was calculated in previous steps
```